



UNIVERSITY OF LEEDS

Machine Learning Introduction

Ali Gooya


a.gooya@leeds.ac.uk



Why bother about Machine Learning?


Machine Learning Job Vacancy Trend

Job postings citing Machine Learning as a proportion of all IT jobs advertised.



Machine Learning Salary Histogram

Salary distribution for jobs citing Machine Learning over the 6 months to 27 September 2020.



What is Machine Learning?

Machine Learning (ML):

The field of study that gives computers the ability to learn without being explicitly programmed.

Arthur Lee Samuel was an American pioneer in the field of computer gaming and artificial intelligence. He popularized the term "machine learning" in 1959.



Further resources on definition:

<https://www.youtube.com/watch?v=ukzFI9rgwfU>



How do you use Machine Learning
in your everyday life?


Some Examples



Spam filtering




Face detection



Pedestrian detection



Movie recommendation



Voice recognition



Types of Machine Learning

Supervised:

Task driven
(classification,
regression)

Unsupervised:

Data driven
(clustering)

Reinforcement:

Goal driven: Algorithm
learns to interact with
environment

Regression




UNIVERSITY OF LEEDS

Task: predict a **continuous** dependent variable y (e.g. weight) given an independent variable x (e.g. height)

Dependent variables are also called **target** variables.

x	y
160	61
154	53
187	79
174	70
165	?



Classification

Task: determine the **discrete** variable y (chair/table) given x (image)



Chair



Chair



Chair



Table



Table



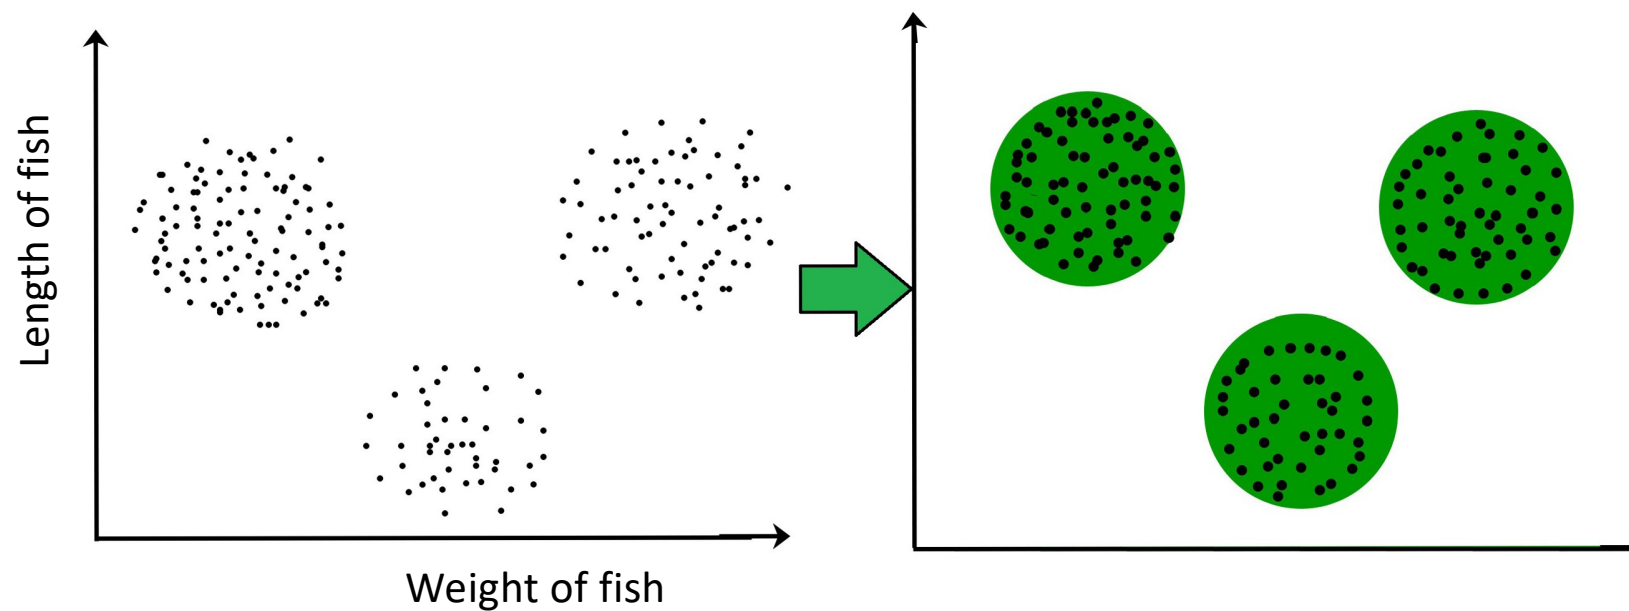
?

Clustering



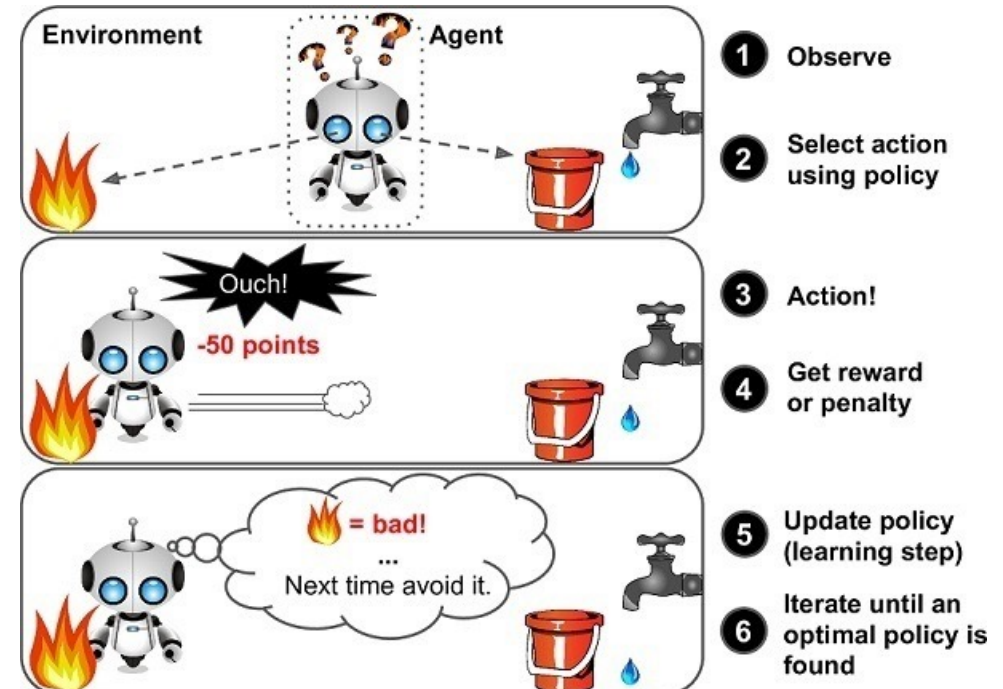

UNIVERSITY OF LEEDS

Task: cluster the data (e.g. height/weight) into coherent groups.



Reinforcement learning

Goal: At any state S_t , the agent learns to take the best action a_t to maximize the reward r_t .



Goals



UNIVERSITY OF LEEDS

- Know the foundation behind the most popular learning algorithms
- Code and apply those algorithms to datasets.

Machine Learning Process




UNIVERSITY OF LEEDS

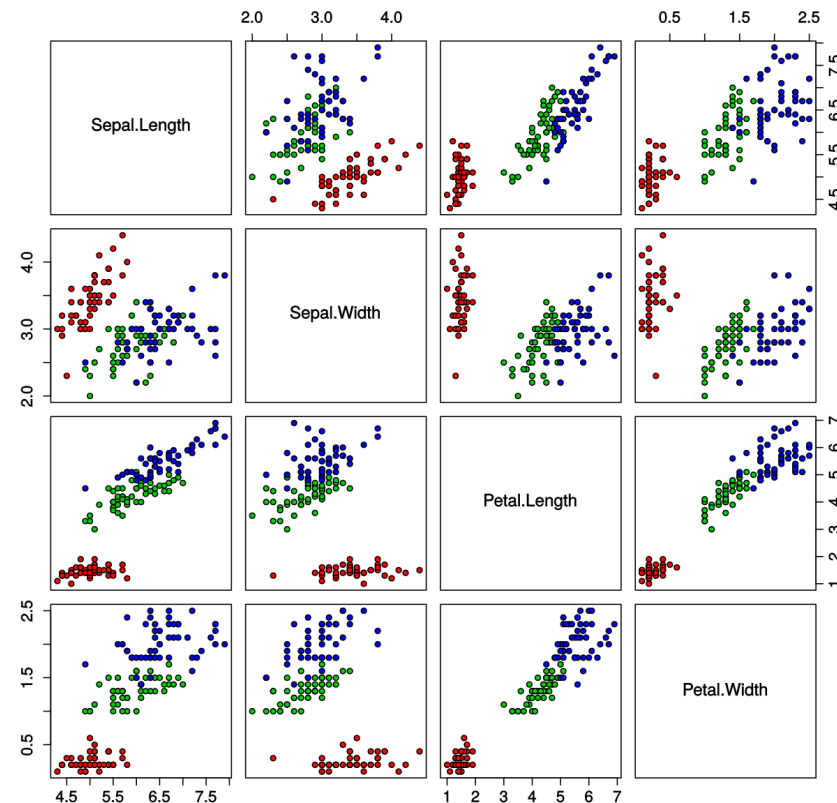
Data Collection:
For example, many images
of various iris flowers



Feature Selection:
E.g., length/widths of petal
and sepal in each flower






Iris Data (red=setosa,green=versicolor,blue=virginica)



Machine Learning Process

Model choice



Training



Data



Tuned
Parameters

Evaluation

		Predicted Class		
		Setosa	Versicolor	Virginica
Actual Class	Setosa	14	1	1
	Versicolor	1	11	3
	Verginica	1	3	10

Machine Learning Process

Data Collection

Feature Selection

Model choice

Training

Evaluation

Most of this class



Other information



UNIVERSITY OF LEEDS

- Lectures: online on Zoom, recordings will be available on Minerva
- Tutorials: will follow the lecture, in-person, every other week (2, 4, ..) unless otherwise said.
- Labs: drop-in, DEC10
- Support: Use Teams page, PGR students
- Assessment plan (tentative):
 - Programming project - Python (20%)
 - In-course online test - Gradescope (30%)
 - Final online test - Gradescope (50%)
- Text books:
 - Machine Learning, An Algorithmic Perspective, Stephen Marsland,
 - Pattern Recognition and Machine Learning, Christopher Bishop